

2-day32-multiple-linear-regression

October 29, 2023

****Multiple Linear Regression**** By:Loga Aswin

```
[26]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
[2]: data = pd.read_csv('/content/winequality-red.csv')

print(data.head())
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[31]: data.tail()
```

```
[31]:      fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
1594           6.2           0.600           0.08           2.0         0.090
1595           5.9           0.550           0.10           2.2         0.062
```

1596	6.3	0.510	0.13	2.3	0.076
1597	5.9	0.645	0.12	2.0	0.075
1598	6.0	0.310	0.47	3.6	0.067

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol	quality
1594	10.5	5
1595	11.2	6
1596	11.0	6
1597	10.2	5
1598	11.0	6

```
[4]: data.describe()
```

```
[4]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806
std	1.741096	0.179060	0.194801	1.409928
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.390000	0.090000	1.900000
50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000

75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[5]: data.isnull().sum()
```

```
[5]: fixed acidity      0
     volatile acidity   0
     citric acid        0
     residual sugar     0
     chlorides          0
     free sulfur dioxide 0
     total sulfur dioxide 0
     density            0
     pH                0
     sulphates          0
     alcohol            0
     quality            0
     dtype: int64
```

Prepare Data:

```
[10]: X = data[['fixed acidity', 'volatile acidity', 'alcohol']].values
      y = data['quality'].values
```

Split Data into Training and Testing Sets:

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=0)
```

Build and Train the Multiple Linear Regression Model:

```
[19]: regressor = LinearRegression()
      regressor.fit(X_train, y_train)
```

```
[19]: LinearRegression()
```

Make Predictions:


```
[20]: y_pred = regressor.predict(X_test)
```

```
[22]: df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
      df1 = df.head(10)
      df1
```

```
[22]:   Actual  Predicted
      0      6    5.940980
      1      5    4.981045
```

2	7	6.417707
3	6	5.422189
4	5	5.886753
5	6	5.063754
6	5	5.263226
7	6	6.068822
8	4	4.823850
9	5	4.962100

```
[28]: df1.plot(kind='bar',figsize=(16,10))
plt.show()
```



Evaluate the Model:

```
[27]: print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
```

Mean Absolute Error: 0.4879795661109293

Mean Squared Error: 0.4096570425100601

Interpret the Results:

```
[30]: coefficients = regressor.coef_
intercept = regressor.intercept_
```

```
print("Coefficients:", coefficients)
print("Intercept:", intercept)
```

Coefficients: [0.03543676 -1.34814989 0.32700295]
Intercept: 2.6602616401275987